

REPLICATED DATABASE SYSTEM - DETAILED FUNCTIONALITY DOCUMENT

PAGE 1: CORE SYSTEM AND PRIMARY OPERATIONS

1. SYSTEM ARCHITECTURE

The system consists of 10 database sites managing 20 shared variables with the following key properties:

- Sites operate independently, can fail/recover separately
- Variables x_1 through x_{20} are maintained across sites
- Even-numbered variables are replicated on all sites
- Odd-numbered variables exist on one site ($\text{site} = i \bmod 10$)
- Each variable x_i has initial value $10i$

2. CORE COMPONENTS AND PRIMARY FUNCTIONALITY

A. Transaction Manager (Main Controller)

Core Functions:

- Transaction Lifecycle:
 - * Begins transactions with unique IDs and timestamps
 - * Executes reads/writes while maintaining consistency
 - * Manages commit/abort decisions
 - * Maintains transaction state and history
- Concurrency Control:
 - * Implements First-Committer-Wins protocol
 - * Detects and prevents cycles in dependency graph
 - * Manages read/write sets for conflict detection
 - * Coordinates between multiple sites
- Operation Handling:
 - * Routes reads to available sites
 - * Manages write propagation
 - * Handles site failures during operations
 - * Queues operations when sites are unavailable

B. Site Manager (Distribution Controller)

Primary Functions:

- Site Status:
 - * Tracks up/down status of each site
 - * Manages site failures and recoveries
 - * Updates accessibility maps
 - * Controls data availability
- Data Distribution:

- * Maps variables to appropriate sites
- * Handles replication rules
- * Manages data consistency across sites
- * Controls access during recovery

C. Data Manager (Storage Controller)

Key Functions:

- Data Storage:

- * Maintains versioned data copies
- * Manages commit logs
- * Provides consistent snapshots
- * Tracks read/write operations

- Version Control:

- * Maintains timestamp-based versions
- * Ensures consistent reads
- * Manages commit history
- * Handles recovery data

3. OPERATION PROCESSING

A. Read Operations ($R(T_i, x_j)$):

1. Validation Steps:

- Check transaction status (not aborted)
- Verify site availability
- Ensure data accessibility
- Validate read permission

2. Execution Process:

- Select appropriate site
- Retrieve consistent version
- Update read sets
- Record operation history
- Return value to transaction

B. Write Operations ($W(T_i, x_j, v)$):

1. Pre-write Checks:

- Validate transaction state
- Check site availability
- Verify write permissions
- Prepare replication targets

2. Write Process:

- Buffer write in transaction

- Update write sets
- Track site dependencies
- Prepare for commit
- Handle replication rules

C. Commit Processing (end(Ti)):

1. Validation Phase:

- Check First-Committer-Wins rule
- Verify site availability
- Detect dependency cycles
- Validate read/write sets

2. Execution Phase:

- Apply buffered writes
- Update commit logs
- Modify site data
- Release resources
- Update transaction status

4. FAILURE AND RECOVERY MANAGEMENT

A. Site Failure Handling (fail(site)):

1. Immediate Actions:

- Mark site as down
- Block access to variables
- Queue pending operations
- Update accessibility maps

2. Transaction Impact:

- Abort affected transactions
- Reassign operations if possible
- Update dependency tracking
- Maintain consistency guarantees

B. Site Recovery Process (recover(site)):

1. Recovery Steps:

- Mark site as available
- Process queued operations
- Update site accessibility
- Restore data access

2. Data Handling:

- Non-replicated: Immediate access
- Replicated: Block reads until write
- Process pending operations
- Restore consistency

5. CONCURRENCY CONTROL IMPLEMENTATION

A. First-Committer-Wins Protocol:

Implementation Details:

- Tracks transaction timestamps
- Maintains write sets
- Compares commit times
- Handles conflicts
- Ensures serializability
- Manages dependencies

B. Dependency Graph Management:

Operation:

- Tracks read/write dependencies
- Maintains transaction ordering
- Detects potential cycles
- Prevents deadlocks
- Manages concurrent access
- Ensures consistency

6. SPECIFIC ALGORITHMS

A. Cycle Detection Algorithm:

Implementation:

- Uses depth-first search
- Maintains visited nodes
- Tracks recursion stack
- Identifies cycles
- Triggers appropriate aborts
- Ensures serializability

B. Site Selection Process:

Steps:

- Evaluates variable location
- Checks site availability
- Considers replication
- Handles failures
- Manages recovery
- Ensures accessibility

7. ERROR SCENARIOS AND HANDLING

A. Common Error Cases:

1. All Sites Down:

- Queue operations
- Maintain consistency
- Handle recovery
- Resume processing

2. Partial Failures:

- Continue operations
- Manage accessibility
- Handle replication
- Maintain consistency

3. Recovery Conflicts:

- Resolve inconsistencies
- Process queued operations
- Restore accessibility
- Update status

B. Transaction Failures:

Handling:

- Clean state restoration
- Resource release
- Dependency updates
- Queue management
- Recovery processing

8. COMMAND PROCESSING AND OUTPUT

A. Input Commands:

Processing:

- begin(T_i): Start transaction
- R(T_i, x_j): Read variable
- W(T_i, x_j, val): Write variable
- end(T_i): End transaction
- fail(site): Site failure
- recover(site): Site recovery
- dump(): System status

B. Output Generation:

Types:

- Transaction status
- Read values
- Commit/abort messages
- Site status updates
- Error notifications
- System state dumps

9. PERFORMANCE CONSIDERATIONS

Key Aspects:

- Minimize blocking
- Optimize site selection
- Efficient recovery
- Quick conflict detection
- Smart queuing
- Effective replication

10. CONSISTENCY GUARANTEES

Ensures:

- Serializable execution
- Consistent reads
- Atomic writes
- Failure atomicity
- Recovery consistency
- Data durability

Classes Information :-

A. Transaction Class :-

```
class Transaction {
    int transaction_id;      // Unique identifier
    int start_timestamp;     // Transaction start time
    int commit_timestamp;    // Commit time if successful
    vector<operation> operations; // List of operations
    bool aborted;           // Transaction status
    vector<int> data_vals;   // Written values (size 20, INT_MAX = no write)
    vector<int> read_vars;   // Variables read by transaction

    // Methods
    void addOperation(operation op);
```

```

    bool canCommit();
    void abort();
    void commit(int timestamp);
}

```

B. Operation Class :-

```

class operation {
    int transaction_id; // Transaction performing operation
    int timestamp;      // Operation timestamp
    char op_type;       // R/W/B/E/F/r (read/write/begin/end/fail/recover)
    int variable_id;    // Target variable (0-19 for x1-x20)
    int site_id;        // Target site (0-9)
    int val;            // Value for writes
    set<int> wait_site_ids; // Sites transaction is waiting for
}

```

C. Transaction Manager :-

```

class TransactionManager {
    // State
    vector<operation> all_operations;
    unordered_map<int, Transaction> transactions;
    vector<set<int>> rw_graph;
    int timestamp;
    set<int> committed_nodes;
    vector<operation> wait_operations;

    // Core Operations
    void beginTransaction(int transaction_id, int timestamp);
    int executeRead(int timestamp_no, int transaction_id, int var_id);
    void executeWrite(int timestamp_no, int transaction_id, int var_id, int val);
    void executeEnd(int timestamp_no, int transaction_id);

    // Validation & Control
    bool can_commit(int transaction_id);
    bool dfs(); // Cycle detection
    bool iscycle(int i, vector<int>&visited, vector<int>&recstack);

    // Dependency Management
    void updateDependencyGraph(Transaction* trans);
    void cleanupTransactionDependencies(int trans_id);
}

```

D. Site Manager :-

```
class SiteManager {
    // State
    vector<bool> site_status;           // Up/down status
    vector<vector<int>>> site_data;      // Current values
    vector<vector<int>>> site_var_map;    // Variable location
    vector<vector<int>>> var_site_map;    // Site mapping
    vector<vector<bool>>> can_read;      // Read accessibility
    vector<vector<int>>> trans_site_write; // Transaction writes

    // Core Operations
    void failsite(int siteno);
    void recover_site(int site_id);
    bool checksite(int siteno);

    // Data Access
    int getsite(int variable_id, int transaction_id);
    void update_all_sites(int var_id, int val, int transaction_id);
    void write_all_sites(int var_id, int val);

    // Initialization
    void initialize();
}
```

E. Data Manager :-

```
class DataManager {
    // State
    map<int, vector<vector<int>>>> main_data; // Timestamped data
    vector<operation> commit_logs;           // Operation history
    vector<int> committed_timestamps;        // Commit times
    vector<set<int>>> read_trans, write_trans; // Transaction sets

    // Core Operations
    int readval(int timestamp, int site_id, int var_id);
    void append_commit(vector<operation>& ops, int timestamp,
                       vector<vector<int>>>& current_data);

    // Version Management
    void cleanup_old_versions();
    void maintain_versions();
}
```


}